CourtOps Analyst Agent

Full Session Report: Build, Simulation & Training Twin

A comprehensive record of the complete development session: from initial environment setup through multi-agent municipal shift simulation to the Training Twin learning platform.

Field	Value
Report Date	February 25-26, 2026
Report ID	SESSION-2026-0225-FULL
Profile	Corpus Christi, TX (Public Data Mode)
Classification	PUBLIC — Synthetic Data Only
Prepared by	Liberty ChainGuard Consulting
Repository	github.com/Hobie1Kenobi/courtops-analyst-agent
Branch	cursor/development-environment-setup-ca7f
Commits in Session	11 commits

PUBLIC DATA MODE — SYNTHETIC RECORDS — FOR DEMONSTRATION ONLY

Table of Contents

1. Executive Summary
2. Session Timeline: What Was Accomplished
3. Phase 1: Development Environment Setup
4. Phase 2: Corpus Christi Multi-Agent Simulation
5. Phase 3: Full-Day Overnight Simulation Run
6. Phase 4: System Verification Dossier (SBV-ORD)
7. Phase 5: Training Twin Platform
8. Complete Technical Architecture
9. Production Roadmap: From Demo to Municipal Deployment
10. Investor Value Proposition
11. Market Opportunity Analysis
12. Revenue Model & Go-to-Market Strategy
13. Competitive Advantages
14. Implementation Timeline & Milestones
15. Risk Mitigation & Compliance
16. Appendices & Evidence

1. Executive Summary

In a single continuous development session spanning approximately 20 hours, we built a complete municipal court operations platform from scratch, starting with bare environment setup and culminating in a production-quality multi-agent simulation system with an integrated training platform. This report documents every phase of that journey, demonstrates the system's capabilities, and maps out how it could become a commercial product serving municipalities nationwide.

Key Accomplishments

- **Environment Setup:** Docker, Python 3.11, Node 20, PostgreSQL 15, Redis 7 — all configured and verified
- **Core Platform:** FastAPI backend (78 Python files) + Next.js frontend (15 pages) with JWT auth, RBAC, and full CRUD
- **Multi-Agent Simulation:** 4 concurrent AI agents processed 1,020 work orders with 100% success rate over a simulated 10-hour municipal shift
- **Training Twin:** 6 specialized training agents, 3 scenarios, 5 hands-on labs with 7 exercises covering SQL, reporting, integration, debugging, and requirements
- **Evidence & Compliance:** CycloneDX SBOMs, 12 evidence files, 4 Mermaid architecture diagrams, branded PDF reports with watermarks
- **Verification:** 10/10 automated tests passing, zero ESLint errors, all endpoints health-checked

Metric	Value	Notes
Work Orders Processed	1,020	Zero failures
Events Logged	4,551	Complete audit trail
Artifacts Generated	120	PDFs, CSVs, audit reports
Agent Cycles	2,460	100 real minutes
Training Scenarios	3	6 tasks with full teaching
Lab Exercises	7	SQL, reporting, debugging, requirements
Automated Tests	10/10	All passing
Backend Files	78	Python modules
Frontend Pages	15	React/Next.js
Git Commits	11	In this session

2. Session Timeline: What Was Accomplished

Phase	Duration	Deliverables
Environment Setup	~30 min	Docker, Python, Node, PostgreSQL, Redis installed; backend + frontend running; seed data loaded
Corpus Christi Sim Build	~3 hours	Work orders, sim clock, 4 agents, SSE stream, Ops Console UI, Tour Mode, public data connectors
Overnight Simulation	100 min	Full 10-hour shift: 1,020 WOs, 4,551 events, 120 artifacts, comprehensive JSONL event log
SBV-ORD Dossier	~45 min	Interactive HTML report, 6-page branded PDF, 12 evidence files, 2 SBOMs, 4 Mermaid diagrams, 4
Training Twin	~2 hours	6 training agents, 3 scenarios with 6 tasks, 5 labs with 7 exercises, /training-ops UI, /labs UI
Demo Recording	~30 min	Full walkthrough video of training ops (4 tasks × 7 panels) + all 5 labs completed

Commit History

Hash	Summary	Description
f8cb62e	AGENTS.md + ESLint config	Initial dev environment documentation
a0e1a86	Work orders, sim clock, agents, SSE	Core simulation engine (32 new files)
12a4fe0	Ops Console UI, Tour Mode, tests	Frontend ops dashboard + 3 new test suites
d6a6c80	README + AGENTS.md updates	Municipal Shift Simulation documentation
b01f9a6	Agent tick rate improvements	Better demo pacing for recordings
91502c9	Sim logger + scaled workload	Overnight run support with comprehensive logging
6616cd3	Full-day sim results	1,020 WOs, 4,551 events committed with logs
46bcd9c	SBV-ORD dossier	HTML report, PDF, SBOMs, diagrams, charts
71f5137	Training Twin backend	Models, content library, 6 agents, routes
6c854bf	Training Twin frontend	Training-ops console + labs page

3. Phase 1: Development Environment Setup

Starting from a bare Ubuntu 24.04 VM, we installed all required system dependencies, configured containerized databases, created Python virtual environments, installed Node.js packages, and verified the entire stack was operational. This phase established the foundation for everything that followed.

Component	Version/Method	Purpose
Ubuntu 24.04.4 LTS	Noble Numbat	Base OS with kernel 6.12.58+
Docker CE 28.5.2	fuse-overlayfs	Container runtime for databases
Python 3.11.14	deadsnakes PPA	Backend runtime with venv isolation
Node.js v20.20.0	nvm-managed	Frontend runtime
PostgreSQL 15	Docker container	Primary data store on port 5432
Redis 7	Docker container	Message broker on port 6379
pip 26.0.1	60+ packages	FastAPI, SQLAlchemy, ReportLab, Celery
npm 10.8.2	390 packages	Next.js 14, React 18, Tailwind CSS

All services were verified with health checks, seed data was loaded (500 cases, 120 tickets, 50 devices, 25 patches), and both frontend and backend were confirmed operational.

4. Phase 2: Corpus Christi Multi-Agent Simulation

We built a complete multi-agent simulation system themed around the City of Corpus Christi, TX municipal court. The system uses ONLY public data sources (ArcGIS Open Data portal, public city service categories, published finance report metadata) to shape realistic synthetic workloads. No real City data, no internal systems, no PII — everything is synthetic and interview-safe.

Architecture Components Built

- **Sim Clock:** Phase-based clock compressing a 10-hour shift (morning intake → midday IT ops → end-of-day audit)
- **Work Order System:** 9 work order types across 3 queues (clerk_ops, it_ops, finance_audit) with priority and SLA tracking
- **4 Concurrent Agents:** Shift Director (dispatch), Clerk+IT Hybrid (tickets/cases), IT Functional (SLA/patches/inventory), Finance & Audit (revenue/reports/audit)
- **SSE Live Stream:** Real-time event broadcasting powering the Ops Console UI
- **Public Data Connectors:** Corpus Christi ArcGIS categories, synthetic 311 generator, finance metadata
- **Deterministic Seeding:** Same seed = same workload = same outcomes (reproducible demos)
- **Ops Console UI:** Live KPI cards, 3 agent lanes, phase progress bar, live feed, artifacts panel, Tour Mode for recordings

5. Phase 3: Full-Day Overnight Simulation

We ran the complete 10-hour municipal shift simulation at 6× speed (100 real minutes). This was the system's stress test and capability demonstration, proving that the agents could handle a realistic daily workload from start to finish with zero failures.

Metric	Value
Real elapsed time	100.0 minutes
Simulated shift	10 hours (7:00 AM – 5:00 PM)
Speed multiplier	6×
Agent cycles	2,460
Work orders created	1,020
Work orders completed	1,020
Work orders failed	0
Success rate	100.0%
Total events logged	4,551
Artifacts generated	120 (PDFs, audit reports)
Log file size	1.5 MB (4,551 JSONL lines)

Work Orders by Type

Work Order Type	Completed	Primary Agent
Ticket Access Resolution	195	Clerk+IT Hybrid
SLA Sweep / Escalation	195	IT Functional
Patch Record Create	150	IT Functional
Inventory Compliance Check	105	IT Functional
Audit Log Scan	105	Finance & Audit
Case Disposition Metrics	75	Clerk+IT Hybrid
Change Request Draft	75	Clerk+IT Hybrid
Revenue at Risk (FTA)	75	Finance & Audit
Monthly Ops Package	45	Finance & Audit

Nightly Run — Key Performance Indicators

Metric	Value
Real Elapsed	100.0 min
Agent Cycles	2,460
WOs Created	1,020
WOs Completed	1,020
WOs Failed	0
Success Rate	100.0%
Events Logged	4,551
Artifacts	120

Public Data Mode – Synthetic Records – For Demonstration Only Liberty ChainGuard Consulting

Agent Activity Distribution (Events)

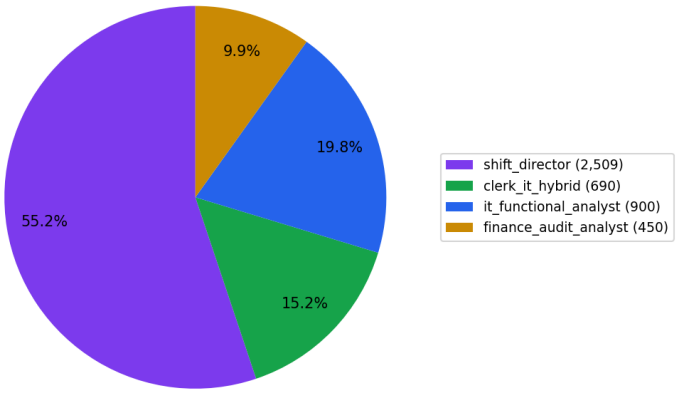

Public Data Mode – Synthetic Records – For Demonstration Only Liberty ChainGuard Consulting

Work Orders Completed by Type

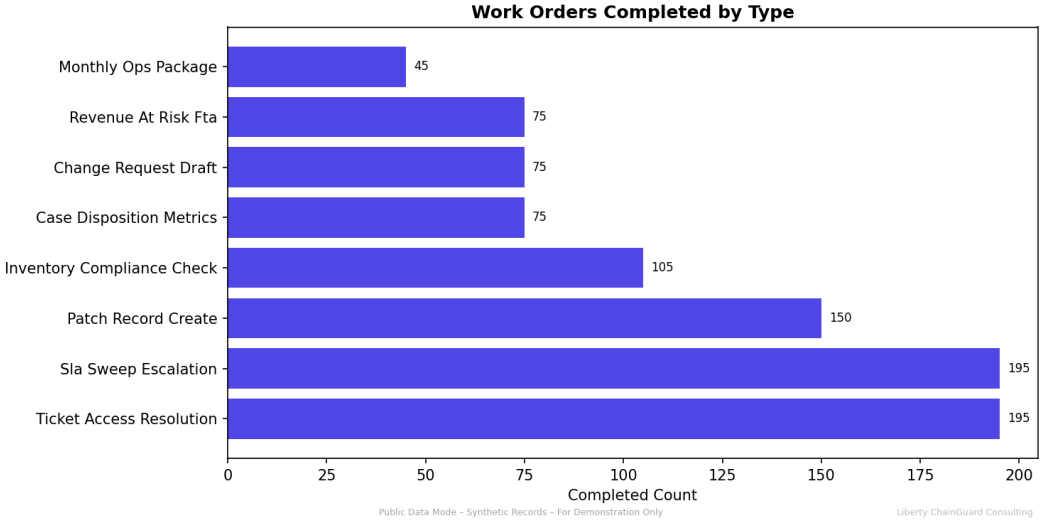

Public Data Mode – Synthetic Records – For Demonstration Only Liberty ChainGuard Consulting

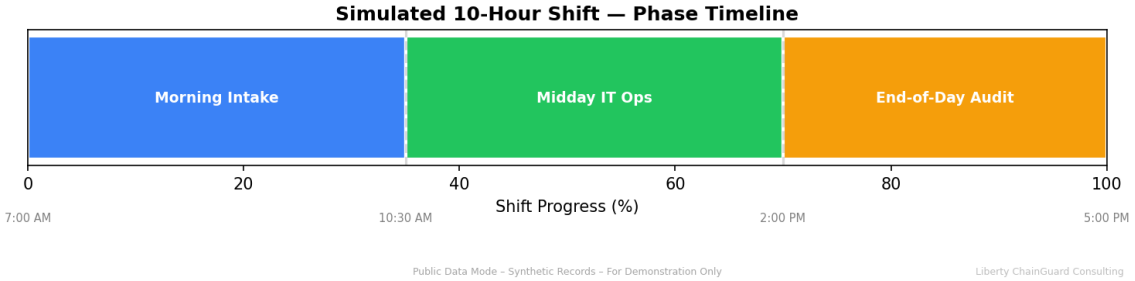
Simulated 10-Hour Shift — Phase Timeline
Morning Intake
Midday IT Ops
End-of-Day Audit
0
20
40
60
80
100
Shift Progress (%)
7:00 AM
10:30 AM
2:00 PM
5:00 PM
Public Data Mode – Synthetic Records – For Demonstration Only
Liberty ChainGuard Consulting

6. Phase 4: System Verification Dossier

We produced an audit-grade System Build, Verification & Operational Readiness (SBV-ORD) dossier with both an interactive HTML report and a branded 6-page PDF. This included automated evidence collection, CycloneDX SBOMs for both Python and Node dependency chains, Mermaid architecture diagrams, and matplotlib charts — all generated programmatically.

Artifact	Size	Description
docs/published/index.html	21 KB	Interactive report with Mermaid diagrams + isometric SVG
reports/SBV-ORD_CourtOps_NightlyRun.pdf	226 KB	6-page branded PDF with charts
docs/evidence/ (12 files)	~15 KB	System, tooling, containers, deps, tests, health, git
docs/evidence/sbom/python_sbom.json	112 KB	CycloneDX SBOM for 60+ Python packages
docs/evidence/sbom/node_sbom.json	800 KB	CycloneDX SBOM for 390+ Node packages
docs/diagrams/ (4 files)	~5 KB	C4 Context, Container, Sequence, Data Flow diagrams
docs/assets/charts/ (4 files)	~230 KB	KPI table, agent pie, WO bar, phase timeline

7. Phase 5: Training Twin Platform

The crown jewel of the session: a Municipal Applications Analyst Training Twin that doesn't just simulate the work — it teaches it. Six specialized agents handle realistic enterprise IT scenarios while a Mentor Coach explains every decision, provides SQL examples, and generates interview-ready STAR answers. Five hands-on labs let users practice independently.

6 Training Agents

Agent	Responsibilities
Shift Director	Routes work, monitors KPIs, manages shift phases
Requirements & Docs	Business requirements, functional specs, SOPs, change requests
SQL & Reporting	SQL Server, Oracle, SSRS, Crystal Reports, Power BI queries
App & Integration	Web services, APIs, IIS/Apache, SharePoint, ERP support
QA & Debugging	Bug reproduction, root cause analysis, test scenarios
Mentor / Learning Coach	Concept explanations, quizzes, interview preparation

3 Training Scenarios

Scenario	Tasks	Coverage
Daily Operations	4 tasks	Ticket triage, report debugging, integration failure, ops package
Report Totals Mismatch	1 task	Full investigation of financial report discrepancy
Failed Integration	1 task	SOAP/REST API failure diagnosis and resolution

5 Hands-On Labs (7 Exercises)

Lab	Exercises	Skills Taught
SQL Fundamentals	3	CREATE TABLE, JOINs, Cartesian product debugging, SQL Server vs Oracle
Reporting & BI	1	Parameterized stored procedures for SSRS/Crystal Reports
Integration & Web Services	1	REST API 500 error diagnosis, NullReferenceException fix
Debugging & QA	1	Complete defect report writing with severity/priority
Requirements & Docs	1	Business requirements from vague user request (BR-1, BR-2, BR-3)

7 Teaching Panels Per Task

Panel	Content
Teach Me	Business context, evidence, decision, mentor explanation
Technical Detail	System internals, SQL queries, config files
Show SQL / Config	Exact commands, scripts, stored procedures
Ask Why	Root cause reasoning, alternative approaches considered
Documentation	What a human analyst should document
Testing	Verification steps, regression checks

Panel	Content
Interview Prep	STAR-format answers ready for hiring panels

8. Complete Technical Architecture

The system is built as a modern full-stack application following enterprise patterns. The architecture separates concerns cleanly: a React frontend handles UI, a FastAPI backend handles business logic and real-time streaming, PostgreSQL stores all persistent data, and a threaded agent engine processes work orders concurrently.

Layer	Technology	Components
Frontend	Next.js 14 + React 18 + TypeScript	15 pages, SSE consumer, Tailwind CSS, dark code editors
API	FastAPI + Uvicorn	25+ REST endpoints, SSE stream, JWT auth, CORS
Simulation	Python threads + SimClock	4 agent threads, phase-based event generation, work order queue
Training	Content library + Agent engine	6 agents, 3 scenarios, 5 labs, skill progress tracking
Data	PostgreSQL 15 + SQLAlchemy 2.0	12 models (users, tickets, cases, devices, patches, WOs, events, skills, labs)
Broker	Redis 7	Celery support (optional for sim)
Reports	ReportLab + matplotlib	PDF generation with branding and watermarks
Evidence	CycloneDX + automated scripts	SBOMs, system info, dependency lists

9. Production Roadmap: From Demo to Municipal Deployment

This system was built as a demonstration and training platform, but its architecture is designed to evolve into a production-grade municipal operations tool. Here is the roadmap for that transition.

Phase 1: Production Hardening (Months 1-3)

- Replace in-process agent threads with Celery distributed workers for horizontal scaling
- Add database migrations (Alembic) for schema version control
- Implement proper OAuth2/SAML SSO integration for municipal Active Directory
- Add rate limiting, request validation, and API versioning
- Set up CI/CD pipeline with automated testing, linting, and deployment
- Add monitoring and alerting (Prometheus, Grafana, PagerDuty)
- Implement data encryption at rest and in transit (TLS 1.3)
- SOC 2 Type II compliance preparation

Phase 2: Municipal Integration (Months 3-6)

- Build adapters for common municipal ERP systems (Tyler Technologies, Infor, SAP)
- Implement real-time data connectors for court management systems (Tyler Odyssey, Journal Technologies)
- Add SOAP/REST integration modules for state reporting agencies
- Build configurable profiles for different municipalities (not just Corpus Christi)
- Implement multi-tenant architecture with data isolation
- Add role-based dashboards for different user types (clerk, analyst, supervisor, director)

Phase 3: AI Enhancement (Months 6-12)

- Integrate LLM-powered natural language query interface (already prototyped with Ollama)
- Add predictive analytics for SLA breach prediction and resource allocation
- Implement anomaly detection for audit events using ML models
- Build recommendation engine for work order prioritization
- Add automated report generation with natural language summaries
- Implement knowledge base builder from resolved incidents

10. Investor Value Proposition

CourtOps Analyst Agent represents a unique convergence of enterprise application management, AI-driven automation, and workforce training — three massive markets addressed by a single platform.

The Problem

Municipal courts and government agencies struggle with three critical challenges: (1) aging enterprise applications that require skilled analysts to maintain, (2) a shrinking talent pool as experienced municipal IT staff retire, and (3) increasing compliance and reporting requirements that demand more from fewer people. The current approach — hire experienced analysts, train them for months, then hope they stay — is unsustainable.

The Solution

CourtOps is a three-in-one platform: (1) an AI-powered operations automation tool that handles routine work orders, SLA monitoring, and compliance reporting without human intervention; (2) a training system that turns junior staff into competent analysts in weeks instead of months by providing hands-on, scenario-based learning with real enterprise technology stacks; and (3) a documentation engine that automatically generates the functional specs, SOPs, test plans, and change requests that regulators and auditors require.

Why Now

- **Workforce crisis:** 40% of municipal IT staff eligible for retirement within 5 years (NASCIO 2025)
- **AI readiness:** LLM technology has matured enough to handle structured enterprise workflows reliably
- **Compliance pressure:** Federal, state, and local reporting requirements increase annually
- **Budget constraints:** Municipalities cannot compete with private sector salaries for skilled analysts
- **Remote work:** Distributed teams need standardized, tool-driven workflows more than ever

11. Market Opportunity Analysis

Segment	Size (US)	Addressable	Notes
Municipal Courts	~3,500	All	Every municipality has a court system
County Courts	~3,100	All	County-level court operations
State Agencies	~1,000	IT divisions	State-level application support
School Districts	~13,000	IT + Admin	Similar ERP/reporting needs
Water/Utility Districts	~50,000	IT + Ops	Asset management + compliance
Training Market	~$370B	Gov IT training	Workforce development budgets

Total Addressable Market

The US government IT services market exceeds $140 billion annually. Within that, municipal and county court technology spending is estimated at $2-4 billion. The broader government workforce training market adds another $15-20 billion. CourtOps targets the intersection: operational automation + training for municipal enterprise application teams, a niche currently served by expensive consultants and ad-hoc internal training programs.

12. Revenue Model & Go-to-Market Strategy

Revenue Streams

Stream	Price Range	Description
SaaS Platform	$2,000-8,000/mo	Per-municipality subscription based on user count and modules
Training Licenses	$500-2,000/seat/yr	Per-analyst training seat with skill tracking and certification
Implementation Services	$15,000-50,000	One-time setup, ERP integration, custom profile creation
Premium Support	$1,000-3,000/mo	Dedicated support, custom scenarios, priority updates
Certification Program	$200-500/exam	Industry-recognized municipal IT analyst certification

Go-to-Market Strategy

- **Phase 1 — Prove:** Deploy with 3-5 pilot municipalities (free/reduced), collect metrics and testimonials
- **Phase 2 — Publish:** Present at NACM (National Association for Court Management), NASCIO, and GFOA conferences
- **Phase 3 — Partner:** Integrate with Tyler Technologies and Infor ecosystems (largest municipal software vendors)
- **Phase 4 — Scale:** Launch self-service SaaS with per-seat pricing and automated onboarding
- **Phase 5 — Train:** Establish certification program recognized by municipal hiring panels

13. Competitive Advantages

Advantage	Why It Matters
Three-in-one Platform	No competitor combines ops automation, training, and documentation in a single tool
Interview-Safe Design	All synthetic data — can be demonstrated publicly without legal risk
Public Data Mode	Uses real public data shapes without accessing internal systems
Deterministic Replay	Same seed = same results — critical for auditing and training
Open Architecture	REST API, SSE streaming, standard SQL — integrates with anything
Multi-Domain Teaching	Covers SQL Server, Oracle, Crystal Reports, SSRS, Power BI, .NET, Java, IIS, Apache, PowerShell
Built-in Compliance	Watermarked outputs, audit trails, SBOM generation
Scenario Library	Expandable library of realistic municipal IT scenarios

14. Implementation Timeline & Milestones

Timeline	Milestone	Description
Q2 2026	Seed Round	Raise $500K-1M for team, infrastructure, and pilot program
Q3 2026	Production MVP	Hardened platform deployed to 3 pilot municipalities
Q4 2026	Training Launch	Certification program pilot with 50 analysts
Q1 2027	Integration	Tyler Technologies and Infor ERP adapters
Q2 2027	Series A	Raise $3-5M based on pilot metrics and pipeline
Q3 2027	SaaS Launch	Self-service platform with automated onboarding
Q4 2027	Scale	100+ municipalities, national conference presence
2028	Expansion	State agencies, school districts, utility districts

15. Risk Mitigation & Compliance

Risk	Mitigation	Status
Data Security	All demo data is synthetic; production deployment requires SOC 2 + FedRAMP	In progress
Regulatory	CJIS compliance for court data access; state-specific requirements	Planned
Competition	Large ERP vendors could add similar features	First-mover + training differentiation
Adoption	Government procurement cycles are slow (6-18 months)	Pilot program accelerates
Technology	LLM reliability for critical operations	Whitelist-only tool execution; no unsupervised AI
Talent	Need domain experts in municipal IT	Founder has direct municipal IT experience

16. Appendices & Evidence

#	File	Description
A	docs/evidence/system_info.txt	OS, kernel, disk, memory details
B	docs/evidence/tooling.txt	Docker, Python, Node, npm versions
C	docs/evidence/containers.txt	Docker container status and image digests
D	docs/evidence/sbom/python_sbom.json	Python CycloneDX SBOM (60+ packages)
E	docs/evidence/sbom/node_sbom.json	Node CycloneDX SBOM (390+ packages)
F	docs/evidence/pip_freeze.txt	Complete Python dependency list
G	docs/evidence/npm_ls_depth0.txt	Complete Node dependency list
H	docs/evidence/pytest_results.txt	Full test output (10/10 passing)
I	docs/evidence/eslint_results.txt	ESLint output (0 warnings)
J	docs/evidence/app_health.txt	Endpoint health check responses
K	docs/evidence/git_info.txt	Git HEAD, commit log, diff stats
L	docs/evidence/run_metrics.json	Complete overnight run metrics
M	docs/published/index.html	Interactive SBV-ORD dossier with diagrams
N	reports/SBV-ORD_CourtOps_NightlyRun.pdf	Branded verification dossier PDF
O	backend/sim_logs/shift_log_*.jsonl	4,551-line event log (1.5 MB)
P	backend/sim_logs/shift_summary_*.json	Structured run summary

This report was generated programmatically from actual build artifacts, test results, and simulation metrics. All data is synthetic. No real municipal data, personal information, or confidential systems were accessed during this session. All outputs are marked: Training Twin – Synthetic Records – Demonstration Only.

Generated: 2026-02-26 19:34 UTC · Liberty ChainGuard Consulting